

M0601 – Project 3

LISC

Learning Instruction Semantics from Code Generators

093311 – Alberto Arruda de Oliveira

Overview

- Binary analysis and instrumentation
 - Program monitoring (Valgrind, Pin)
 - Virtualization
 - Malware Analysis
 - Etc.
- Modeling of instruction Semantics
 - Translation of assembly to intermediate language representation
- Requires previously existing models, created manually
 - Limited Support of new architectures

LISC

Objective

- Automated modeling of instruction semantics
- Compilers: Intermediate Language -> Assembly Code
 - GCC, LLVM, etc.
- LISC:
 - Learn from code generators
 - Automatically extract semantics
 - Test extracted semantics

LISC

Benefits

- Automated instruction semantics modeling
 - Machine Learning Approach
 - Reduce manual efforts
 - Broaden Architecture Support
- Architecture Neutrality
 - Code Generators
- Well tested compiler code
 - Among the most tested code

Project 3

Objective

- Test LISC in a already used architecture of choice
 - x86, ARM, AVR
- Reproduce **at least one** of LISC's evaluation approaches:
 - Completeness
 - Architecture-Neutrality
 - Performance
 - Correctness

LISC Evaluation

Completeness

- Are all the instructions lifted correctly?
- 99.5% of Ubuntu/x86 and 99.8% of Debian/ARM binaries
 - Missed Instructions are mostly NOPS
- Evaluation:
 - Testing programs P_{test} and training programs $P_{train} \subset P_{test}$
 - Build a transducer using P_{train}
 - Translate the assembly of P_{test}
 - Metrics:
 - Fraction of unique instructions translated
 - Total number of non-translated instructions

LISC Evaluation

Architecture Neutrality

- Can the system be easily adapted to a new architecture?
- 9 person-hours to support ARM and 3 person-hours to support AVR
- **Evaluation:**
 - **Code that is architecture specific**
- **Project 4:** Support a new architecture!

LISC Evaluation Performance

- How long does training/testing takes?
- 10 minutes training
 - 1.4GB of code generator logs
 - 4830 transducer states
- Lifting binaries from Ubuntu/x86 and Debian/ARM architectures (around 100M instructions): around 8 hours
- Linear-time algorithm for lifting binaries

LISC Evaluation

Correctness

- Semantic equivalence test
 - *ArCheck*
- Consistency Check
 - Distinct Intermediate Language Snippets -> Same Assembly Instruction A
 - A may translate to two different snippets
- Loopback Test: Given a binary with a list of instructions α
 - Lift α to Intermediate language γ using LISC
 - Use GCC to generate code α' for γ
 - Check if α' is consistent with α

Project 4

Possible Ideas

- New check of architecture neutrality
 - Generate language model for a non-supported architecture
- Alternative machine learning technique
 - Can a different approach to training and testing be employed?
- Evaluation Flaws
 - Is there some aspect of current evaluation that fails to cover possible scenarios?

THANK YOU!